

Phase 0 Prep — Scope of Work for Sameer & Gnana

LocalAI TV · What the Architecture Team Needs · Prep Only (Gate Held)

To: Sheikh Sameer & Gnana Rajnan (AI Pipeline Engineers) **From:** LocalAI TV Architecture Team **For:** Founder Koneti Mohan Reddy (to forward) **Date:** 16 May 2026 **Status:** 🟡 **Prep only.** Phase 0 execution stays held until the Founder lifts the gate. This document defines the four prep tasks and exactly what to send back.

0. SECURITY RULE — READ FIRST

This is **not** about trust. It is standard practice and it protects LocalAI TV.

Topic	Rule
Sameer holding root / all internal credentials	✅ Fine — the Founder's decision. Normal for a lead engineer. Not in question.
Real secrets (passwords, API keys, access keys, SSH keys)	Live only in the shared team vault (Bitwarden) among Founder + Sameer + Gnana .
What the Architecture Team receives	Non-secret status text ONLY (Section 3 template). Never a credential.
What is never sent to the Architecture Team — by ANY channel	No credentials via chat, file upload, email, WhatsApp, screen-share, or AnyDesk . There is no exception and no need for one.
AnyDesk / remote desktop / screen-share to transfer anything to the Architecture Team	❌ Not used. Not needed. Do not set this up.

Why: an AI assistant and any chat transcript are not a vault. The Architecture Team produces plans, code, and instructions; **you** operate the live systems with the real credentials. The two never meet. Trusting Sameer fully does not change this — the boundary is about *where secrets live*, not *who is trusted*.

Net: keep every secret in the vault. Send the Architecture Team only the plain-text status in Section 3.

1. The Four Prep Tasks (Part A — no gate crossed)

#	Task	Owner	Notes
A0	Submit the AWS GPU quota increase — today	Sameer/Gnana	AWS Console → Service Quotas → EC2 → "Running On-Demand G and VT instances" → region ap-south-1 → request ≥ 8 vCPUs . Free to request; AWS approval takes 1–2 business days — this is the long pole, start it first.
A1	Send the non-secret status confirmations (Section 3)	Sameer/Gnana	Plain text only. ~15–30 min. This is the one item the Architecture Team needs to verify Phase-0 readiness.
A2	Confirm the shared vault	Founder + Sameer/Gnana	Bitwarden (free): one shared "LocalAI TV" vault; Sameer & Gnana invited; credentials placed there by the team. Never sent to the Architecture Team.
A3	AWS guard-rails on the account	Sameer/Gnana	Dedicated least-privilege IAM user (only S3 + the Phase 0 EC2 instance), separate bucket localaitv-content-mumbai (ap-south-1), billing budget alert (~₹35,000/mo), tag all Phase 0 resources project=phase0 — so testing cannot touch the live channels.

2. Timeline (since the team already has all access)

Item	Team effort	External wait	Realistic completion
A0	~15 min to submit	AWS approval 1–2 business days (out of our hands)	Submit today → cleared in 1–2 days
A1	~15–30 min	none	Same day — within the hour
A2	~1–2 h (one-time vault setup)	none	Same day
A3	~2–4 h	none	Same day (scope the IAM policy carefully)

Bottom line: if started today, prep is effectively complete in **~1 working day**, bounded only by AWS's GPU-quota approval (1–2 days). Nothing here touches the gate.

3. A1 — Send Exactly This (fill in, plain text, NO secrets)

Phase 0 Readiness — Status (non-secret)

AWS

- Account ready: Yes / No
- Region: ap-south-1 (Mumbai)? Yes / No
- Same AWS account as the currently-live channels? Yes / No
- GPU quota increase (A0) submitted? Yes / No — date: _____

Database

- PostgreSQL version: _____ (expected: 15)
- Running where: Hostinger VPS / other: _____
- Reachable from where the pipeline runs? Yes / No

Domain / DNS

- `localaitv.com` DNS managed at which registrar: _____
- Is `content.localaitv.com` free to point at the CDN? Yes / No

VPS

- Hostinger OS + version: _____
- Pipeline currently runs here? Yes / No

That is the entire content the Architecture Team needs. **No usernames, passwords, keys, or screenshots of secrets — ever.**

4. Gate Status (unchanged)

- This is **prep only**. Phase 0 **execution** (launch EC2, swap to NVENC, wire S3, measure G1–G4 — see PHASE-0-RUNBOOK Part B) does **not** start until the **Founder explicitly lifts the gate** and scope is frozen (architecture-review clarifications C1–C4 still pending with the AI team).
- Doing A0–A3 now does not cross the gate — it removes start-day delay.

Sign-off

Party	Action	Status
Sameer & Gnana	Do A0, A2, A3; send A1 status (Section 3)	☐
Founder — Koneti Mohan Reddy	Vault access for the team; lift the gate when ready (separate step)	☐
Architecture Team	Verify readiness from A1; flag gaps	☐ Awaiting A1

LocalAI TV Architecture Team · LocalAI Media Network Pvt Ltd · CIN: U63910KA2025PTC212593 · Hyderabad, India

Prep only. No credentials to the Architecture Team by any channel. Plan v1.3 unchanged. Execution waits for the Founder's gate-lift.